

DA_HW2

Howard

2025-03-05

Problem 1

Consider the simple linear regression model

$$y = c + \beta_1 x + \epsilon,$$

where the intercept c is known.

(a) Derivation of the Least-Squares Estimator of β_1

Minimize the sum of squared errors:

$$RSS(\beta_1) = \sum_{i=1}^n \left(y_i - c - \beta_1 x_i \right)^2.$$

Differentiate $RSS(\beta_1)$ with respect to β_1 and set the derivative to zero:

$$\frac{\partial RSS}{\partial \beta_1} = -2 \sum_{i=1}^n x_i (y_i - c - \beta_1 x_i) = 0.$$

This implies:

$$\sum_{i=1}^n x_i (y_i - c) - \beta_1 \sum_{i=1}^n x_i^2 = 0.$$

Solving for β_1 , we obtain:

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n x_i (y_i - c)}{\sum_{i=1}^n x_i^2}.$$

This estimator is reasonable because it essentially regresses the centered response $y_i - c$ on x_i , thereby isolating the effect of x on y .

(b) Variance of $\hat{\beta}_1$

Given the model $y_i = c + \beta_1 x_i + \epsilon_i$ and noting that

$$\hat{\beta}_1 = \beta_1 + \frac{\sum_{i=1}^n x_i \epsilon_i}{\sum_{i=1}^n x_i^2},$$

and assuming that ϵ_i are independent with $\text{Var}(\epsilon_i) = \sigma^2$, the variance of $\hat{\beta}_1$ is:

$$\text{Var}(\hat{\beta}_1) = \frac{\sigma^2}{\sum_{i=1}^n x_i^2}.$$

(c) Confidence Interval for β_1

Assuming the errors $\epsilon_i \sim N(0, \sigma^2)$, the estimator $\hat{\beta}_1$ is normally distributed. However, since σ^2 is usually unknown, we estimate it by:

$$\hat{\sigma}^2 = \frac{1}{n-1} \sum_{i=1}^n \left(y_i - c - \hat{\beta}_1 x_i \right)^2,$$

where the degrees of freedom is $n - 1$ (because only β_1 is estimated).

Then the statistic

$$t = \frac{\hat{\beta}_1 - \beta_1}{\sqrt{\hat{\sigma}^2 / \sum_{i=1}^n x_i^2}}$$

follows a t -distribution with $n - 1$ degrees of freedom. Thus, a $100(1 - \alpha)\%$ confidence interval for β_1 is given by:

$$\hat{\beta}_1 \pm t_{n-1, \alpha/2} \sqrt{\frac{\hat{\sigma}^2}{\sum_{i=1}^n x_i^2}},$$

where $t_{n-1, \alpha/2}$ is the critical value from the t -distribution with $n - 1$ degrees of freedom.

Comparison with the Case When the Intercept Is Unknown

When the intercept is unknown, the least-squares estimator of β_1 is

$$\tilde{\beta}_1 = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2},$$

with variance

$$\text{Var}(\tilde{\beta}_1) = \frac{\sigma^2}{\sum_{i=1}^n (x_i - \bar{x})^2}.$$

Because

$$\sum_{i=1}^n (x_i - \bar{x})^2 = \sum_{i=1}^n x_i^2 - n\bar{x}^2,$$

this sum is smaller than $\sum_{i=1}^n x_i^2$ (unless $\bar{x} = 0$). Consequently, the variance—and hence the confidence interval—is narrower when the intercept c is known.

Problem 2

(a) Convert the 400 images into a 400 × 2576 data matrix. Add an additional column indicating the physical gender label.

List and Read PNG Images

```
library(png)
image_dir <- "D:/DA_HW/DAHW/ORL Faces"
image_files <- list.files(image_dir, pattern = "\\\\.png$", full.names = TRUE)
n <- length(image_files)
if(n != 400){
  stop("Expected 400 images, but found ", n)
}
```

Convert Images to a Data Matrix

```
data_matrix <- matrix(NA, nrow = n, ncol = 46 * 56)
for (i in 1:n) {
  img <- readPNG(image_files[i])
  data_matrix[i, ] <- as.vector(t(img))
}
dim(data_matrix)
```

```
## [1] 400 2576
```

Append the Gender Label

- **Classify the g by manual.**

```
# There are 40 subjects, each with 10 images.
subject_ids <- rep(1:40, each = 10)

# male (1) and female (0).
subject_gender <- c(0,rep(1,6),0,1,0,1,0,rep(1,19),0,rep(1,8))

# Create a gender vector for all 400 images:
gender_labels <- subject_gender[subject_ids]

# Append the gender labels as an additional column to the data matrix.
final_data <- cbind(data_matrix, gender = gender_labels)

# Check the dimensions of the final data matrix.
dim(final_data)
```

```
## [1] 400 2577
```

(b)Regress the gender label on all the 2576 pixels? What do you observe?

Fit a Logistic Regression on All Pixels

- Because the predicted valued is binary,So I thinking using Logistic Regression is the proper regress method.

```
# Convert final_data (matrix) to a data frame
df <- as.data.frame(final_data)
#names
pixel_names <- paste0("pixel", 1:2576)
colnames(df) <- c(pixel_names, "gender")
#Fit the Logistic Regression for binary response(gender = 0/1)
model_glm <- glm(gender ~ ., data = df,family = binomial)
s <- summary(model_glm)
print(s$coefficients[1:200])
```

```
## [1] 159.36374 781.67373 -847.97814 4132.16870 -1505.34419 78.55472
## [7] -2335.17627 467.06981 -944.42390 -1327.05245 828.05710 220.75294
## [13] 55.92137 -603.85901 369.55916 -509.05263 981.80508 -434.72787
## [19] 735.32352 23.75581 762.45535 -513.81574 227.33318 -1100.63479
## [25] 1328.37120 -1063.43422 1836.92143 -1077.17562 368.28560 -762.24699
## [31] 772.16267 -360.81218 -1323.68696 1408.91931 -585.32422 1328.85971
## [37] -1263.23467 2663.00645 340.72642 -743.59963 1885.03681 3493.06216
## [43] -502.27881 -27.99947 -1763.74385 -4162.54316 -2895.92094 143.93773
## [49] -2175.22095 3485.58077 -2095.45336 500.75854 53.79110 1270.80878
## [55] 411.84152 1196.93241 -1838.86515 451.48558 -348.56685 1739.46126
## [61] -1161.34255 -28.16144 -838.39268 314.46349 1296.34321 -1659.48584
## [67] -1417.32145 1040.81201 -533.16037 776.09260 800.26049 -2347.81683
## [73] 2108.39340 -438.25105 -533.38410 -595.66634 -550.78434 -963.81107
## [79] 597.94529 1765.45249 49.45760 -1522.55324 836.91697 -3491.14019
## [85] 586.50565 -3117.07163 602.04702 -2637.38396 -1487.82409 -3737.64655
## [91] 4424.48404 2955.70501 3022.72058 -3462.76600 -1687.81616 -762.62803
## [97] 93.09846 257.62831 -225.02492 836.00995 -2434.72318 2049.22764
## [103] -625.89656 176.58230 -1168.79181 83.71876 705.07207 -1016.71674
## [109] 1886.62843 -2245.08138 1391.25305 1947.52269 -1445.76777 -353.74296
## [115] 1239.45675 547.53479 -1889.73058 1533.50759 -720.98089 47.36489
## [121] 139.75813 -2951.32977 3526.81847 -1133.71428 1444.74560 -27.30775
## [127] 151.27809 -743.32376 -965.43718 2283.04942 1196.75206 63.50886
## [133] 2086.54104 -1524.22715 -1628.34090 903.73755 -1032.29360 549.30811
## [139] 3119.47475 -1867.81772 4032.06622 -245.71497 716.20895 -1865.63878
## [145] 805.68048 -897.32426 2198.53437 -1089.22153 645.74200 71.36172
## [151] 2612.50959 -2154.30703 2075.38177 -1697.00650 1039.19271 -618.84286
## [157] 659.66269 -2251.38484 1968.81497 -603.99859 730.14379 -2451.62835
## [163] 1078.50836 -319.38788 244.53508 818.12681 1860.23544 -2370.00334
## [169] 1879.95634 -710.69152 -496.29313 419.32721 150.27125 -1405.36643
## [175] -651.86554 754.03169 -2238.81007 284.15360 -760.59163 -1681.53257
## [181] 4157.63040 -377.30888 1433.21036 165.01893 -5530.84316 5824.24475
## [187] -1577.16195 876.19625 559.72526 1478.47693 -1922.88391 822.30552
## [193] -281.86754 -374.10463 -1145.15256 -1472.60154 -242.80459 1748.03985
## [199] -2973.69879 2442.65283
```

What do we observe

- **Convergence Problems:** With 2576 features for only 400 samples, the logistic regression might fail to converge or produce perfect separations.
- **Many Coefficient with Large Standard Errors:** Because the model is extremely high dimensional, so many parameters are unstable.
- **$n \ll p$.** With 2576 predictors and only 400 samples, the model is severely overparameterized. Many predictors are highly collinear, so many coefficients are set to NA.
- **in my opinion** This regression highlights the need for dimensionality reduction (like PCA), or regularization methods (like, Ridge, Lasso) when working with high dimensional image data.

(c) Perform the stepwise regression from a null model to find the important pixels. Plot the chosen pixels on a 46×56 canvas.

Stepwise Feature Selection and Pixel Plot

```
library(MASS) # for stepAIC
library(stringr)
```

```
## Warning: 套件 'stringr' 是用 R 版本 4.2.3 來建造的
```

```
df$gender <- as.factor(df$gender) # logistic regression expects a factor
```

Define a Null and Full Model for Logistic Regression

```
# Null model: includes only the intercept
null_model <- glm(gender ~ 1, data = df, family = binomial)
```

```
# Full model: includes all pixels
full_model <- glm(gender ~ ., data = df, family = binomial)
```

```
## Warning: glm.fit:演算法沒有聚合
```

Forward Stepwise Selection (Up to 50 Steps)

```
stepwise_model <- stepAIC(
  object = null_model,
  scope = list(lower = null_model, upper = full_model),
  direction = "forward",
  trace = FALSE, # set to TRUE to see iteration details
  steps = 50      # limit the maximum number of forward steps
)

summary(stepwise_model)
```

```
##
## Call:
## glm(formula = gender ~ pixel2320 + pixel685 + pixel1558 + pixel607 +
##      pixel1701 + pixel2574 + pixel1475 + pixel2534 + pixel208 +
##      pixel452 + pixel1401 + pixel1878 + pixel2332 + pixel1049,
##      family = binomial, data = df)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -5.678e-04  2.000e-08  2.000e-08  2.000e-08  5.080e-04
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    3958.4    166857.5   0.024   0.981
## pixel2320     -3242.5    136932.9  -0.024   0.981
## pixel685      -3236.0    136104.6  -0.024   0.981
## pixel1558      2371.0    102844.8   0.023   0.982
## pixel607        947.7     44840.1   0.021   0.983
## pixel1701      1835.2     83277.9   0.022   0.982
## pixel2574     -2100.5     90737.2  -0.023   0.982
## pixel1475     -1578.0     69222.6  -0.023   0.982
## pixel2534      1392.0     60491.8   0.023   0.982
## pixel208      -1745.1     80635.0  -0.022   0.983
## pixel452       1766.3     81781.3   0.022   0.983
## pixel1401     -2049.9     91745.4  -0.022   0.982
## pixel1878      1287.7     56104.4   0.023   0.982
## pixel2332     -1484.6     69783.6  -0.021   0.983
## pixel1049      -708.0     54728.7  -0.013   0.990
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 3.0142e+02  on 399  degrees of freedom
## Residual deviance: 1.5914e-06  on 385  degrees of freedom
## AIC: 30
##
## Number of Fisher Scoring iterations: 25
```

Identify the Selected Pixels

```
selected_features <- names(coef(stepwise_model))[-1]
selected_features
```

```
## [1] "pixel2320" "pixel685" "pixel1558" "pixel607" "pixel1701" "pixel2574"
## [7] "pixel1475" "pixel2534" "pixel208" "pixel452" "pixel1401" "pixel1878"
## [13] "pixel2332" "pixel1049"
```

```
pixel_idx <- as.numeric(sub("pixel", "", selected_features))
pixel_idx
```

```
## [1] 2320 685 1558 607 1701 2574 1475 2534 208 452 1401 1878 2332 1049
```

Plot the Important Pixels on a 46×56 Canvas


```
# Initialize all zeros
important_pixel_map <- rep(0, 46 * 56)

# Mark selected pixel indices as 1
important_pixel_map[pixel_idx] <- 1

# Convert to a 46x56 matrix. We assume row-major flattening.
pixel_matrix <- matrix(important_pixel_map, nrow = 46, ncol = 56, byrow = TRUE)

# We'll use base R's image() function with two colors:
image(1:56, 1:46,
      t(apply(pixel_matrix, 2, rev)), # flip to put row=1 at top
      col = c("white", "red"),
      xlab = "Column (1→56)",
      ylab = "Row (1→46)",
      main = "Important Pixels Selected (up to 50)")
```

Important Pixels Selected (up to 50)



```
# 'red' = selected pixel
# 'white' = not selected
```

Problem 3

Step1 Load the Volcano Data

- Using R built-in `volcano` dataset.
- `volcano_data[i,j]` = height at ($x_1 = i, x_2 = j$).
- We want to find the (x_1, x_2) that gives the max height.

```
volcano_data <- volcano
dim(volcano_data)
```

```
## [1] 87 61
```

Step2 Set Initial Parameters and Initialize the Search

- When I set the half-Window size as 5, the Highest Point is (27,37), And the Elevation at Highest Point is 180.
- So I change the `window_size` to ten will find the highest point.
- The half-window size for local regression which refers to how far out from the current point you look in each direction when defining a local neighborhood for regression.

```
# Starting position: (x1, x2) = (87, 1)
current_x1 <- nrow(volcano_data) # 87
current_x2 <- 1
# Record the search path as a List of (x1, x2) coordinates
path <- list(c(current_x1, current_x2))

# Parameters for the search
window_size <- 8 # half-window size for local regression
step_size <- 1 # step size
```

Step3 Iterative Search for the Peak

- *Local Window Definition*: Determines which part of the data to use for local approximation.
- *Grid Creation and Elevation Extraction*: Constructs a coordinate grid with corresponding elevation values.
- *Local Linear Regression*: Fits a plane to the local data and extracts the gradient.
- *Movement*: Uses the gradient and step size to update the search position.
- *Convergence Check*: Stops the algorithm when no further improvement is detected.

```

for (i in 1:100) {
  row_min <- max(1, current_x1 - window_size)
  row_max <- min(nrow(volcano_data), current_x1 + window_size)
  col_min <- max(1, current_x2 - window_size)
  col_max <- min(ncol(volcano_data), current_x2 + window_size)

  # Create a grid of coordinates for the local window
  x1_seq <- row_min:row_max
  x2_seq <- col_min:col_max
  grid <- expand.grid(x1 = x1_seq, x2 = x2_seq)
  # Extract the corresponding elevations.
  # Note: volcano_data[row, col] where row corresponds to x1 and col to x2.
  grid$z <- as.vector(volcano_data[x1_seq, col_min:col_max])

  # Fit a local linear model:  $z \sim x1 + x2$ 
  model <- lm(z ~ x1 + x2, data = grid)
  coefs <- coef(model)
  # Extract the slopes (gradient components)
  grad <- coefs[c("x1", "x2")]

  # Determine move direction: sign(gradient) multiplied by step_size
  move_x1 <- sign(grad["x1"]) * step_size
  move_x2 <- sign(grad["x2"]) * step_size

  # Update position
  new_x1 <- current_x1 + move_x1
  new_x2 <- current_x2 + move_x2

  # Ensure the new position is within bounds
  new_x1 <- max(1, min(new_x1, nrow(volcano_data)))
  new_x2 <- max(1, min(new_x2, ncol(volcano_data)))

  # If the new position is unchanged, assume a local maximum has been reached
  if ((new_x1 == current_x1) && (new_x2 == current_x2)) {
    break
  }

  # Update the current position and add it to the path
  current_x1 <- new_x1
  current_x2 <- new_x2
  path <- c(path, list(c(current_x1, current_x2)))
}

```

Step4 Report the Highest Point

```

highest_point <- c(current_x1, current_x2) # (x1, x2) in 1-based indexing
highest_elevation <- volcano_data[current_x1, current_x2]

cat("Highest Point (1-based index):", highest_point, "\n")

```

```
## Highest Point (1-based index): 21 29
```

```
cat("Elevation at Highest Point:", highest_elevation, "\n")
```

```
## Elevation at Highest Point: 191
```

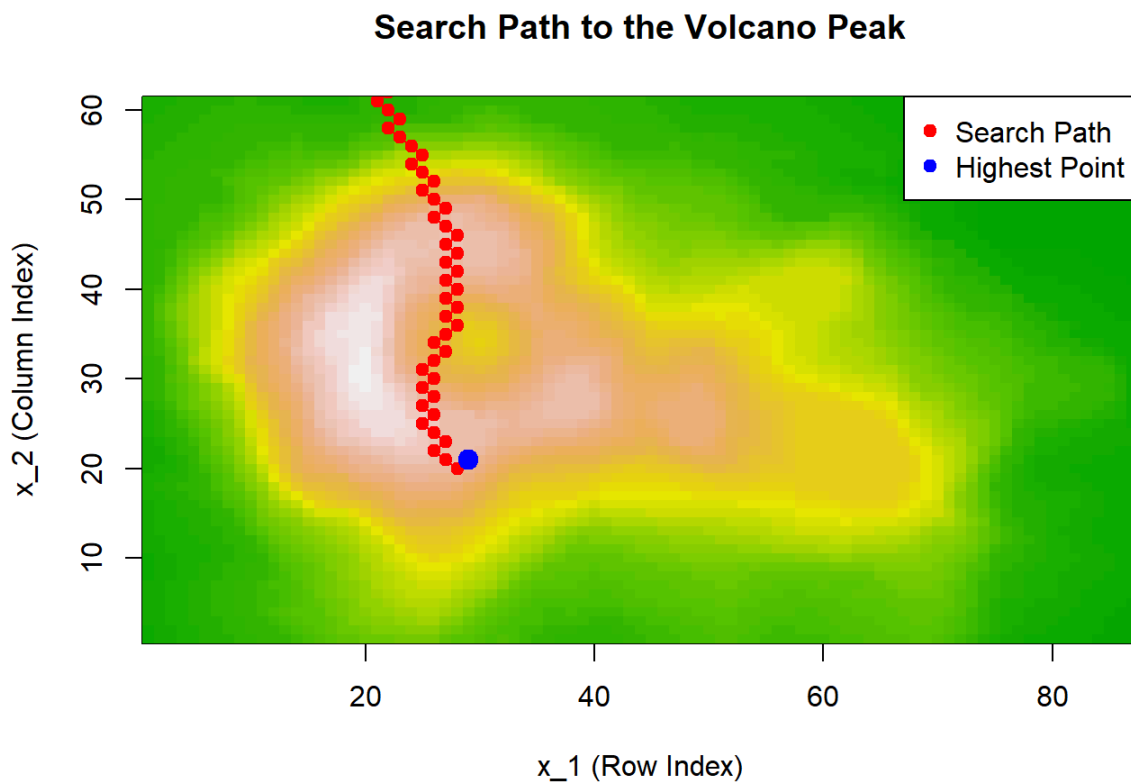
Step5 Visualize the Path on a Contour Plot


```
# Create a grid for plotting: x1 = rows, x2 = columns
x1_range <- 1:nrow(volcano_data)
x2_range <- 1:ncol(volcano_data)

# Use image() to display the volcano data
image( x1_range, x2_range, volcano_data,
      col = terrain.colors(50),
      xlab = "x_1 (Row Index)",
      ylab = "x_2 (Column Index)",
      main = "Search Path to the Volcano Peak")

# Convert the path (list) to a matrix for plotting.
path_mat <- do.call(rbind, path)

# In the plot, x corresponds to columns (x2) and y to rows (x1).
lines(path_mat[,2], path_mat[,1], col = "red", type = "o", pch = 19)
points(highest_point[2], highest_point[1], col = "blue", pch = 19, cex = 1.5)
legend("topright", legend = c("Search Path", "Highest Point"),
      col = c("red", "blue"), pch = c(19, 19))
```



Problem 4

- We assume the model $y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \epsilon$, $\epsilon \sim N(0, \sigma^2)$.

(a) Use a regression package in R to analyze the problem and interpret the results.

- Simulate the data with $x_1 \sim N(0, 1)$, $x_2 \sim N(0, 1)$
- And $y = 2 + 3x_1 + (-1)x_2 + \epsilon$
- where $\epsilon \sim N(0, \sigma^2)$

1. Simulate the Data

```

n <- 50000
beta0 <- 2
beta1 <- 3
beta2 <- -1
sigma <- 2 # standard deviation of the noise

# Generate predictors from N(0,1)
x1 <- rnorm(n, mean = 0, sd = 1)
x2 <- rnorm(n, mean = 0, sd = 1)

# Generate noise term: epsilon ~ N(0, sigma^2)
epsilon <- rnorm(n, mean = 0, sd = sigma)

# Generate response variable y
y <- beta0 + beta1 * x1 + beta2 * x2 + epsilon

# Combine into a data frame
df <- data.frame(y, x1, x2)

# Display the first few rows of the data
head(df)

```

```

##           y           x1           x2
## 1  0.3731012 -0.1126949 -0.2658292
## 2  3.4582951 -0.6940615 -0.9665603
## 3 -0.1095448 -1.4536800 -0.9555962
## 4 -2.0788741 -1.0576313  1.7959559
## 5 -0.1108864 -1.0796584 -0.1700542
## 6  0.7635127 -0.6163644 -0.2222210

```

2. Fit the Linear Model

```

model <- lm(y ~ x1 + x2, data = df)
summary(model)

```

```

##
## Call:
## lm(formula = y ~ x1 + x2, data = df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -8.067 -1.345  0.000  1.336  8.438
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  2.006376   0.008929   224.7  <2e-16 ***
## x1           2.999468   0.008948   335.2  <2e-16 ***
## x2          -1.013448   0.008892  -114.0  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.996 on 49997 degrees of freedom
## Multiple R-squared:  0.7157, Adjusted R-squared:  0.7157
## F-statistic: 6.292e+04 on 2 and 49997 DF,  p-value: < 2.2e-16

```

3. Explanation and Interpretation

- The fitting model is $\hat{y} = 2.011668 + 3.004628x_1 - 1.00265x_2$
- Extremely low p-values ($< 2.2e-16$) for all coefficients indicate that each is statistically significant at typical thresholds (0.05, 0.01, etc.).
- Very large t-values (225.5, 358.1, -111.5) confirm the significance.

- R^2 around 71.5% of the variance in y is explain by x_1, x_2
- Overall Model Significance F-statistic is 6.281×10^4 , p-value < $2.2e-16$, strongly indicating the model is significant and reject the null hypothesis that all slope coefficients are zero.

(b) Estimating Regression Coefficients via Gradient Descent

1. Define the Loss Function and Its Gradient

```
# Loss function (mean squared error) multiplied by 1/2 for convenience.
loss_function <- function(beta, y, x1, x2) {
  n <- length(y)
  predictions <- beta[1] + beta[2]*x1 + beta[3]*x2
  loss <- sum((y - predictions)^2) / (2*n)
  return(loss)
}

# Gradient of the loss function with respect to beta parameters.
gradient_function <- function(beta, y, x1, x2) {
  n <- length(y)
  predictions <- beta[1] + beta[2]*x1 + beta[3]*x2
  error <- y - predictions
  grad0 <- -sum(error) / n
  grad1 <- -sum(error * x1) / n
  grad2 <- -sum(error * x2) / n
  return(c(grad0, grad1, grad2))
}
```

2. Implement Gradient Descent

```
# Initialization
beta <- c(0, 0, 0) # initial guess for (beta0, beta1, beta2)
alpha <- 0.05      # Learning rate (you may adjust this)
max_iter <- 1000
tol <- 1e-6

# To store the loss and parameters at each iteration
loss_history <- numeric(max_iter)
beta_history <- matrix(0, nrow = max_iter, ncol = 3)

for (iter in 1:max_iter) {
  current_loss <- loss_function(beta, y, x1, x2)
  loss_history[iter] <- current_loss
  beta_history[iter, ] <- beta

  grad <- gradient_function(beta, y, x1, x2)
  beta_new <- beta - alpha * grad

  # Check for convergence: change in beta is small
  if (sum(abs(beta_new - beta)) < tol) {
    beta <- beta_new
    beta_history[iter, ] <- beta
    loss_history[iter] <- loss_function(beta, y, x1, x2)
    cat("Converged at iteration", iter, "\n")
    loss_history <- loss_history[1:iter]
    beta_history <- beta_history[1:iter, ]
    break
  }

  beta <- beta_new
}
```

```
## Converged at iteration 247
```

```
cat("Final beta from gradient descent:\n")
```

```
## Final beta from gradient descent:
```

```
print(beta)
```

```
## [1]  2.006369  2.999458 -1.013446
```

```
cat("Final loss:", tail(loss_history, 1), "\n")
```

```
## Final loss: 1.992845
```

3.Compare with OLS Solution

```
ols_model <- lm(y ~ x1 + x2, data = df)
ols_beta <- coef(ols_model)
cat("OLS coefficients:\n")
```

```
## OLS coefficients:
```


```
print(ols_beta)
```

```
## (Intercept)          x1          x2
##    2.006376    2.999468   -1.013448
```

4.Plot the Evolution of the Loss Function and contour plot

```
plot(loss_history, type = "l", col = "blue", lwd = 2,
     xlab = "Iteration", ylab = "Loss (MSE/2)",
     main = "Evolution of the Loss Function")
```

Evolution of the Loss Function



```

# Fix beta0 at its final value
beta0_final <- beta[1]

# Define a grid for beta1 and beta2
beta1_range <- seq(ols_beta[2] - 1, ols_beta[2] + 1, length.out = 100)
beta2_range <- seq(ols_beta[3] - 1, ols_beta[3] + 1, length.out = 100)
grid <- expand.grid(beta1 = beta1_range, beta2 = beta2_range)

# Compute the loss over the grid for each pair (fixing beta0)
grid$loss <- apply(grid, 1, function(b) {
  loss_function(c(beta0_final, b["beta1"], b["beta2"]), y, x1, x2)
})

# Create a contour plot
contour(beta1_range, beta2_range,
        matrix(grid$loss, nrow = 100, byrow = TRUE),
        xlab = expression(beta[1]), ylab = expression(beta[2]),
        main = "Search Path in ( $\beta_1$ ,  $\beta_2$ ) Space ( $\beta_0$  fixed)",
        nlevels = 30, col = "gray")

# Overlay the search path for beta1 and beta2
lines(beta_history[,2], beta_history[,3], col = "red", lwd = 2)
points(beta_history[,2], beta_history[,3], col = "red", pch = 19)

```

Search Path in (β_1 , β_2) Space (β_0 fixed)

